



# HPC Master Trainer Workshop – February 2024

## MPI Collective Communication

Sandeep Agrawal  
[sandeepa@cdac.in](mailto:sandeepa@cdac.in)






# Agenda

- Collective Communication
- Types Of Collective Communication
- Collective Communication Routines
- Collective Communication II Routines
- Some Advance MPI Features

# Point To Point Communication


Simplest form of message communication

- Message is sent from a sending process to a receiving process only these two process need to know anything about the message.



# Collective Communication

- Collective communication must involve all processes in the scope of a communicator.
- Collective Participation in solving the problem.






# Characteristics

MPI collective communication routines differ in many ways from MPI point-to-point communication routines

- Involves coordinated communication within a *group* of processes identified by an MPI communicator.
- Substitute for a more complex sequence of point-to-point calls.
- All routines block until they are locally complete.
- In some cases, a *root* process originates or receives all data.
- Amount of data sent must exactly match amount of data specified by receiver.
- No message tags are needed.

# Communicators and Groups

- MPI uses objects called communicators and groups to define which collection of processes may communicate with each other. Most MPI routines require you to specify a communicator as an argument.





# Collective Communication



- Collective communication must involve all processes in the scope of a communicator. All processes are by default, members in the communicator **MPI\_COMM\_WORLD**
- It is the programmer's responsibility to ensure that all processes within a communicator participate in any collective operations.



# Types of Collective Operations

- **Synchronization** Processes wait until all members of the group have reached the synchronization point.
- **Data Movement** broadcast, scatter/gather, all to all.
- **Collective Computation (reductions)** - one member of the group collects data from the other members and performs an operation (min, max, add, multiply, etc.) on that data.




---

# Collective Communication Routines

# Barrier

- Creates a barrier synchronization in a group.
- Each task, when reaching the MPI\_Barrier call, blocks until all tasks in the group reach the same MPI\_Barrier call.







# Data Movement

MPI provides three types of collective data movement routines

- Broadcast
- Gather
- Scatter


# Broadcast

- Broadcasts (sends) a message from the process with rank "root" to all other processes in the group.



# Broadcast

- Sends data stored in buffer buf of process source to all the other processes in the group comm.





# Broadcast

`int MPI_Bcast(void *buf, int count, MPI_Datatype datatype, int source, MPI_Comm comm)` - C Program

**buf** : the address of the send buffer.

**count** : the number of elements sent to each process.


**datatype** is MPI defined constant indicating the data type of the elements in the buffer.

**root** : is an integer indicating the rank of broadcast

**comm** : the communicator.

- Distributes distinct messages from a single source task to each task in the group. if one wants to distribute the data into  $n$  equal segments, where the  $i^{\text{th}}$  segment is sent to the  $i^{\text{th}}$  process in the group which has  $n$  processes.

```
MPI_Scatter(&sendbuf, sendcnt, sendtype, &recvbuf  
, recvcnt, recvtype, root, comm)
```






# Scatter

**MPI\_Scatter(&sendbuf,sendcnt,sendtype,&recvbuf,recvcnt,recvtype,root,comm)**

- |          |   |   |
|----------|---|---|
| sendbuf  | : | the address of the send buffer.               |
| sendcnt  | : | the number of elements sent to each process.  |
| sendtype | : | the data type of the send buffer elements.    |
| recvbuf  | : | the address of the receive buffer.            |
| recvcnt  | : | the number of elements in the receive buffer. |
| recvtype | : | the data type of the receive buffer elements. |
| root     | : | the rank of the sending process.              |
| comm     | : | the communicator.                             |

## Gather

- Gathers distinct messages from each task in the group to a single destination task. This routine is the reverse operation of MPI\_Scatter.





# Gather

**MPI\_Gather(&sendbuf,sendcnt,sendtype,&recvbuf,recvcount,recvtype,root,comm)**

**sendbuf** : the address of the send buffer.

**sendcnt** : the number of elements in the send

**sendtype** : the data type of the send buffer elements

**recvbuf** : the starting address of the receive buffer.

**recvcnt** : the number of elements for any single receive.

**recvtype** : the data type of the receive buffer elements.


**root** : the rank of the receiving process.

**comm** : the communicator.

# AllGather

- Concatenation of data to all tasks in a group. Each rank in the group in effect performs a one - to - all broadcast.


```
int MPI_Allgather(void *sendbuf, int  
sendcount,MPI_Datatype senddatatype, void*recvbuf,  
int recvcount,MPI_Datatype recvdatatype, MPI_Comm  
comm)
```



# All-to-All


Each process sends a different portion of sendbuf to each other process (incl. itself)

- `recvbuf` of target process stores data in rank order
  - `sendcount` specifies no. of elements sent to each process



# Reduce


- Used to combine partial results from all processors
  - Result returned to root processor
  - Several types of operations available
  - Works on single elements and arrays




# Reduce

**MPI\_Reduce (&sendbuf,&recvbuf,count,datatype,op,root,comm)**

**MPI\_REDUCE** (*sendbuf*,*recvbuf*,*count*,*datatype*,*op*,*root*,*comm*,*ierr*)



MPI SUM





# Reduction operations



Operation	Meaning	Datatypes
MPI_MAX	Maximum	C integers and floating point
MPI_MIN	Minimum	C integers and floating point
MPI_SUM	Sum	C integers and floating point
MPI_PROD	Product	C integers and floating point
MPI_LAND	Logical AND	C integers
MPI_BAND	Bit-wise AND	C integers and byte
...	...	...



# MPI Collective Communication



The gather, scatter, allgather, and alltoall routines have variable data versions. For their variable data versions, each process can send and/or receive a different number of elements.

- MPI\_Scatterv
- MPI\_Gatherv
- MPI\_Allgatherv
- MPI\_Alltoallv

What does the "v" stand for?

varying – size,relative location of the messages.



# Summary

PO	P1	P2*	P3	Function	PO	P1	P2	P3
a	b	c	d	MPI_Gather			a,b,c,d	
a	b	c	d	MPI_Allgather	a,b,c,d	a,b,c,d	a,b,c,d	a,b,c,d
	a,b,c,d			MPI_Scatter	a	b	c	d
a,b,c,d	e,f,g,h	i,j,k,l	m,n,o,p	MPI_AlltoAll	a,e,i,m	b,f,j,n	c,g,k,o	d,h,l,p
	b			MPI_Bcast	b	b	b	b
SBuf	SBuf	SBuf	SBuf		RBuf	RBuf	RBuf	RBuf

- Sender/Root process required by MPI\_Gather, MPI\_Scatter, MPI\_Bcast



# Some Advance MPI Features

- **Dynamic Processes** - extensions that remove the static process model of MPI. Provides routines to create new processes
- **One-Sided Communications** - provides routines for one directional communications. Include shared memory operations (put/get) and remote operations.
- **Extended Collective Operations** - allows for non-blocking collective operations and application of collective operations to inter-communicators
- **Parallel I/O** - describes MPI support for parallel I/O



# Pros and Cons

## Pros of MPI Collective Communication

- Substitute for a more complex sequence of point-to-point calls.
- Communications may, or may not, be synchronized (implementation dependent).
- No message tags are needed.



# Pros and Cons



## Cons of MPI Collective Communication

- A great deal of hidden communication takes place with collective communication.
- Performance depends greatly on the particular implementation of MPI
- There may be forced synchronization, not always best to use collective communication

ધ્ન્યવાદ ધ્ન્યવાદ ધ્ન્યવાદ  
ધ્ન્યવાદ ધ્ન્યવાદ ધ્ન્યવાદ  
ધ્ન્યવાદ ધ્ન્યવાદ ધ્ન્યવાદ  
**Thank You**

